

TECHNICAL BRIEF

Why We Must Automate Our Data Pipelines

THE RISK — When Manual Entry Goes Wrong

In March 2024, a fuel distribution company in East Africa experienced an unplanned shutdown at one of its depots. The cause was traced back to a single spreadsheet. A night-shift supervisor, fatigued at the end of a 12-hour shift, entered a pressure reading of 4.7 as 47 — a single missed decimal point. The figure fed into the morning report without challenge, because nobody checks a number that looks plausible at a glance. By the time the operations team realised the sensor log did not match actual equipment behaviour, four hours of throughput data had been recorded against the wrong baseline. The correction required two engineers, a full audit of that week's logs, and a six-hour recalibration window that cost the depot an estimated 30,000 USD in lost delivery capacity.

This is not an unusual story. It is the predictable outcome of asking human beings to perform the same repetitive data entry task hundreds of times a week without error. Research across manufacturing and logistics consistently shows that manual data entry carries an error rate of 1 to 4 percent. In a plant logging 144 sensor readings per day, that means between 1 and 6 incorrect values entering the system every 24 hours — silently, without any alert.

THE SOLUTION — Automatic Irrigation, Not Carrying Buckets

Automating a data pipeline is exactly like installing an irrigation system in place of carrying water buckets by hand. The water still travels from the source to the field — but it moves on a fixed schedule, through validated channels, without anyone having to remember to do it. If a pipe is blocked or the pressure is wrong, a sensor triggers an alert. Nobody has to be standing there watching.

An automated data pipeline works the same way. The raw sensor readings are collected from the source file, checked against a set of rules — is this pressure value physically possible? is this operator ID formatted correctly? — and loaded into a clean database, every morning at 06:00, without any staff involvement. If a reading fails a quality check, the pipeline stops

and sends an alert. Nothing corrupt reaches the database. Nothing reaches the morning report unless it has already been verified.

THE VALUE — What Automation Delivers

- Time saved. The current process of pulling, cleaning, and checking the weekly sensor log takes an analyst approximately 3 hours. Automated, this takes 4 seconds. Across 52 weeks, that is over 150 analyst-hours per year redirected from data entry to actual analysis.
- Errors eliminated. The pipeline we built this week identified 18 physically impossible sensor readings in a single 7-day log — values that would have passed through a manual process unchallenged. Automated quality checks catch these before they reach any report.
- Faster decisions. When the pipeline runs at 06:00, the operations manager has a verified, clean dataset on their desk before the morning stand-up. Under the current process, that same dataset arrives — if it arrives at all — by mid-morning. A six-hour head start on a pressure anomaly is the difference between a controlled response and an emergency shutdown.

Excel is a powerful tool. It is the right tool for scenario planning, budgeting, and one-off calculations. It is the wrong tool for a process that runs the same steps on the same data every single day. That is exactly what a pipeline is for.